

Import Libraries and Data

```
In [1]: import quandl
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib
import matplotlib.pyplot as plt
import os
import statsmodels.api as sm
import warnings

warnings.filterwarnings("ignore")

plt.style.use('fivethirtyeight')
```

```
In [2]: # Configure API key
quandl.ApiConfig.api_key = 'Zhv1MKMFgFbxUx9fv4qY'
```

```
In [3]: # Importing the Federal Reserve Economic Data\ Civilian Unemployment Rate
data = quandl.get('FRED/UNRATE')
```

```
In [4]: # Checking the head
data.head(5)
```

Out[4]:

	Value
Date	
1948-01-01	3.4
1948-02-01	3.8
1948-03-01	4.0
1948-04-01	3.9
1948-05-01	3.5

```
In [5]: data.shape
```

```
Out[5]: (889, 1)
```

```
In [6]: data.columns
```


```
Out[6]: Index(['Value'], dtype='object')
```

```
In [7]: type(data) # This may not be important to copy to Final draft
```

```
Out[7]: pandas.core.frame.DataFrame
```

```
In [42]: # Plot the data using matplotlib
plt.figure(figsize=(15,5), dpi=100)
plt.title('Civilian Unemployment Rates- (1950 through 2020)'),
plt.xlabel('Years'),
plt.ylabel('Values'),
plt.plot(data)
```

```
Out[42]: [<matplotlib.lines.Line2D at 0x1c8d3827590>]
```



2. Subsetting, wrangling, and cleaning time-series data

```
In [9]: # Resetting index to use Date column as a filter
data_2 = data.reset_index()
```

```
In [10]: data_2.head()
```

Out[10]:

	Date	Value
0	1948-01-01	3.4
1	1948-02-01	3.8
2	1948-03-01	4.0
3	1948-04-01	3.9
4	1948-05-01	3.5

```
In [11]: # Subsetting df
data_sub = data_2.loc[(data_2['Date'] >= '2010-01-01') & (data_2['Date'] < '2020-06-01')]
```

I didn't find a data set that correlated with my final project "Flavors of Cocoa", so I opted to do my time series over a dataset that sparked my interest: "the unemployment rate" amongst civilians over the last ten years (2010- 2020). This analysis will show patterns or trends to help pinpoint a possible cause of the unemployment rate spike(s).

```
In [12]: # Checking shape of subset
data_sub.shape
```

Out[12]: (125, 2)

```
In [13]: data_sub.head()
```

Out[13]:

	Date	Value
744	2010-01-01	9.8
745	2010-02-01	9.8
746	2010-03-01	9.9
747	2010-04-01	9.9
748	2010-05-01	9.6

```
In [14]: # Setting Date column as the Index and creating a datetime column
data_sub['datetime'] = pd.to_datetime(data_sub['Date'])

# Now setting the datetime as the index of the df
data_sub = data_sub.set_index('datetime')

# Lastly, dropping the date column
data_sub.drop(['Date'], axis=1, inplace=True)
```


```
In [15]: # Checking the head after making changes
data_sub.head()
```

Out[15]:

	Value
datetime	
2010-01-01	9.8
2010-02-01	9.8
2010-03-01	9.9
2010-04-01	9.9
2010-05-01	9.6

```
In [37]: # Plot the new data set
plt.figure(figsize=(15,5), dpi=100)
plt.title('Civilian Unemployment Rates- Last 10 years'),
plt.xlabel('Years'),
plt.ylabel('Values'),
plt.plot(data_sub)
```

```
Out[37]: [<matplotlib.lines.Line2D at 0x1c8d33bf090>]
```



Cleaning Time Series Data ##### Before Decomposing

```
In [17]: # Checking for missing values
data_sub.isnull().sum()
```

```
Out[17]: Value      0
dtype: int64
```

```
In [18]: # Check for duplicates
dups = data_sub.duplicated()
```

```
In [19]: dups.sum()
```

```
Out[19]: 71
```

```
In [20]: dups
```

```
Out[20]: datetime
2010-01-01    False
2010-02-01     True
2010-03-01    False
2010-04-01     True
2010-05-01    False
...
2020-01-01     True
2020-02-01     True
2020-03-01     True
2020-04-01    False
2020-05-01    False
Length: 125, dtype: bool
```

Not addressing dups. ## Not true duplicates- not removing. Month and day repeats itself yearly.


3. Time-series analysis: decomposition

```
In [21]: # Decompose the time series using an additive model
decomposition = sm.tsa.seasonal_decompose(data_sub, model='additive')
```

```
In [22]: # Define a fixed size for all special charts.
from pylab import rcParams
```

```
In [23]: rcParams['figure.figsize'] = 18, 7
```

```
In [39]: # Plot the separate components
decomposition.plot(),
plt.xlabel('Years'),
# Show the plot
plt.show()
```



Analyzing the data:

Trend line: After 2010, the trend line is downward, showing a gradual decrease in the unemployment rate, with a spike in 2020 due to the COVID-19 pandemic of 2020-21.

Seasonality line: displays spikey curves throughout 2010-2020, demonstrating some form of seasonality. We can surmise that the spikes are during the holiday season (Christmas).

Residual line: doesn't appear to show much noise. However, the data points start to trail off around 2020, which is unsurprising.

4. Testing for stationarity

```
In [25]: # Import adfuller function
from statsmodels.tsa.stattools import adfuller
```


```
In [26]: # Define the function
def dickey_fuller(timeseries):
    # Perform the Dickey-Fuller test:
    print('Dickey-Fuller Stationarity test:')
    test = adfuller(timeseries, autolag='AIC')
    result = pd.Series(test[0:4], index=['Test Statistic', 'p-value', 'Number of Lags Used', 'Number of Observations Used'])
    for key, value in test[4].items():
        result['Critical Value (%s)' % key] = value
    print(result)
```

```
In [27]: # Apply the test using the function on the time series
dickey_fuller(data_sub['Value'])
```

```
Dickey-Fuller Stationarity test:
Test Statistic          -1.644966
p-value                  0.459667
Number of Lags Used      0.000000
Number of Observations Used  124.000000
Critical Value (1%)      -3.484220
Critical Value (5%)      -2.885145
Critical Value (10%)     -2.579359
dtype: float64
```

In reviewing the test results, my test statistics of -1.644966 is more significant than all of the critical values, implying that my series is non-stationary. Therefore, I must apply further differentiation to reject the hypothesis that the unemployment rate(s) fluctuates significantly yearly.


```
In [28]: # Import the autocorrelation and partial correlation and #Check out a plot of autocorrelations
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
plot_acf(data_sub)
plt.show()
```



5. Stationarizing the Data

```
In [29]: # Using the df.shift() function
data_diff = data_sub - data_sub.shift(1)
```

```
In [30]: data_diff.dropna(inplace = True)
```

```
In [31]: data_diff.head()
```

```
Out[31]:
```


	Value
datetime	
2010-02-01	0.0
2010-03-01	0.1
2010-04-01	0.0
2010-05-01	-0.3
2010-06-01	-0.2

```
In [32]: data_diff.columns
```

```
Out[32]: Index(['Value'], dtype='object')
```

```
In [41]: # Check out what the differencing did to the time-series curve
plt.figure(figsize=(15,5), dpi=100)
plt.plot(data_diff)
# Add an axis labels
plt.title('Data Differencing Time Series Curve'),
plt.xlabel('Years'),
plt.ylabel('Values')
```


```
Out[41]: Text(0, 0.5, 'Values')
```



```
In [34]: dickey_fuller(data_diff)
```

```
Dickey-Fuller Stationarity test:
Test Statistic      -1.154772e+01
p-value             3.507950e-21
Number of Lags Used  0.000000e+00
Number of Observations Used  1.230000e+02
Critical Value (1%)  -3.484667e+00
Critical Value (5%)  -2.885340e+00
Critical Value (10%) -2.579463e+00
dtype: float64
```

```
In [35]: # Checking Autocorrelation
plot_acf(data_diff)
plt.show()
```



```
In [ ]:
```